



CLOUD COMPUTING

Assignment 03



ALI ABBAS
BCSF17A502

```

HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0
Job Counters:
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=2239
  Total time spent by all reduces in occupied slots (ms)=2504
  Total time spent by all map tasks (ms)=2239
  Total time spent by all reduce tasks (ms)=2504
  Total vcore-milliseconds taken by all map tasks=2239
  Total vcore-milliseconds taken by all reduce tasks=2504
  Total megabyte-milliseconds taken by all map tasks=2292736
  Total megabyte-milliseconds taken by all reduce tasks=2564896
Map-Reduce Framework
  Map input records=3
  Map output records=8
  Map output bytes=83
  Map output materialized bytes=80
  Input split bytes=110
  Combine input records=0
  Combine output records=0
  Reduce input groups=6
  Reduce shuffle bytes=80
  Reduce input records=6
  Reduce output records=6
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=154
  CPU time spent (ms)=1450
  Physical memory (bytes) snapshot=524750848
  Virtual memory (bytes) snapshot=5219487744
  Total committed heap usage (bytes)=456130560
  Peak Map Physical memory (bytes)=292489344
  Peak Map Virtual memory (bytes)=2606850048
  Peak Reduce Physical memory (bytes)=232341504
  Peak Reduce Virtual memory (bytes)=2612637696
Shuffle Errors:
  BAD_ID=0
  CONNECTION=0
  ID_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=31

```

Activities | Terminal | Mon 08:01

Cloud-Computing-Proje... | Google | Cloud Computing BS CS... | How to Install Hadoop... | Apache Download Mirror... | +

← → ↻ | https://phoenixnap.com/kb/install-hadoop-ubuntu | ☆ | 📄 | 📌 | ☰

Bare Metal Cloud now available at special prices! [SEE DISCOUNTS](#)

Please only use the backup mirrors to download KVM, P2P signatures and patches (P2P 400) - or if no other mirrors are working.

HTTP
https://downloads.apache.org/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz

```

hadoop@dev-HP-ProBook-450-G7: ~
File Edit View Search Terminal Help
o.B000
. 0.+0=.0.+
00.*=0.0E.+*
-----[SHA256]-----
hadoop@dev-HP-ProBook-450-G7:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
hadoop@dev-HP-ProBook-450-G7:~$ chmod 0600 ~/.ssh/authorized_keys
hadoop@dev-HP-ProBook-450-G7:~$ ssh localhost
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is SHA256:CH8qIMzToFuV++d+05SyPRHjL9dsQTD5MxPoCit8gnA.
Are you sure you want to continue connecting (yes/no)? y
Please type 'yes' or 'no':
Host key verification failed.
hadoop@dev-HP-ProBook-450-G7:~$ wget https://downloads.apache.org/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz
--2021-06-14 07:59:23-- https://downloads.apache.org/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz
Resolving downloads.apache.org (downloads.apache.org)... 88.99.95.219, 135.181.209.10, 135.181.214.104, ...
Connecting to downloads.apache.org (downloads.apache.org)|88.99.95.219|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 359196911 (343M) [application/x-gzip]
Saving to: 'hadoop-3.2.1.tar.gz'

hadoop-3.2.1.tar.gz      22%[=====] 77.04M  1.55MB/s  eta 6m 17s
Resolving downloads.apache.org (downloads.apache.org)... 88.99.95.219, 135.181.209.10, 135.181.214.104, ...
Connecting to downloads.apache.org (downloads.apache.org)|88.99.95.219|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 359196911 (343M) [application/x-gzip]
Saving to: 'hadoop-3.2.1.tar.gz'

hadoop-3.2.1.tar.gz      100%[=====] 342.56M  10.2MB/s   in 31s
2020-04-22 09:29:23 (10.9 MB/s) - 'hadoop-3.2.1.tar.gz' saved [359196911/359196911]

```

error sites. The official

3.1.4.tar.gz

hadoop-3.2.1

default)

hadoop-3.2.1/hadoop-3.2.1/hadoop-3.2

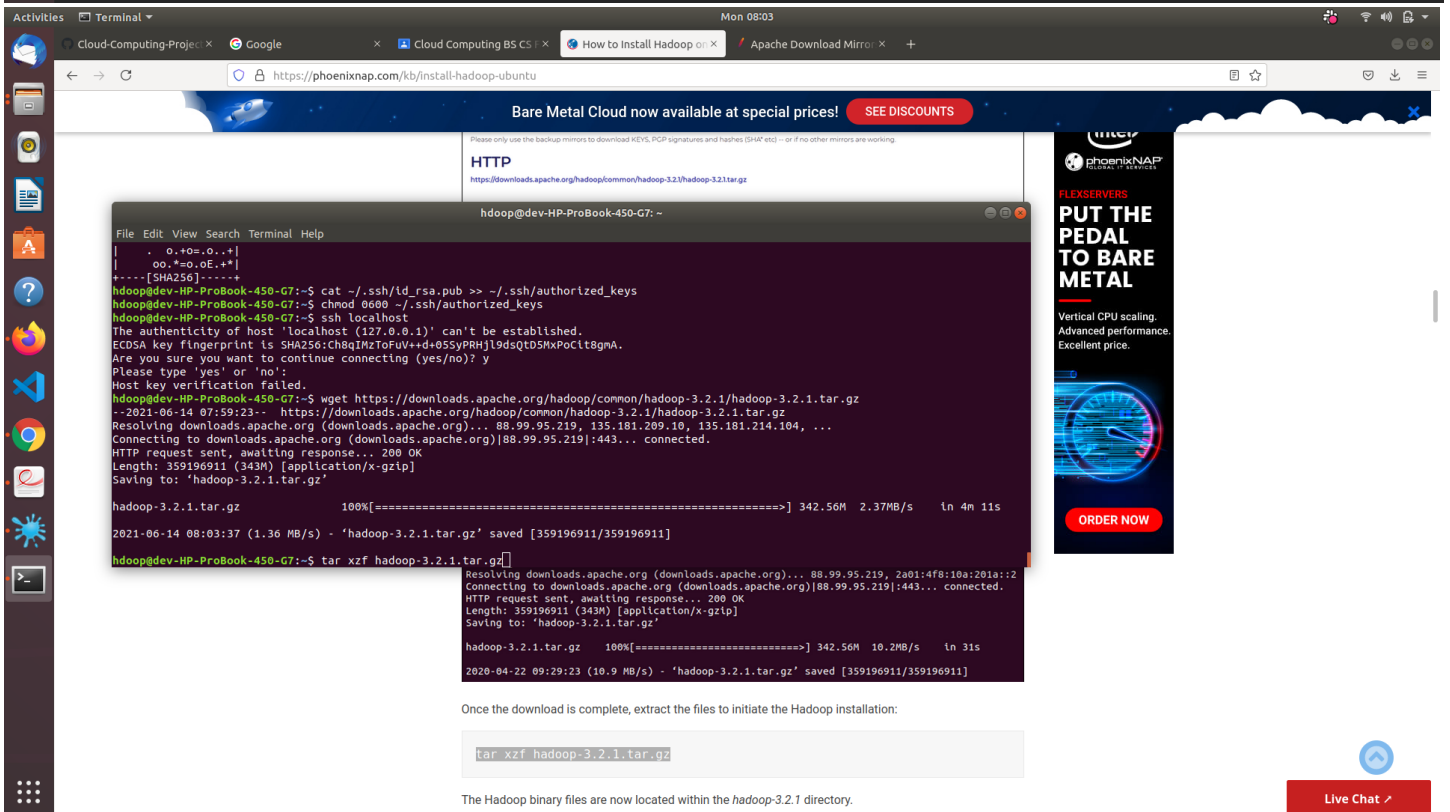
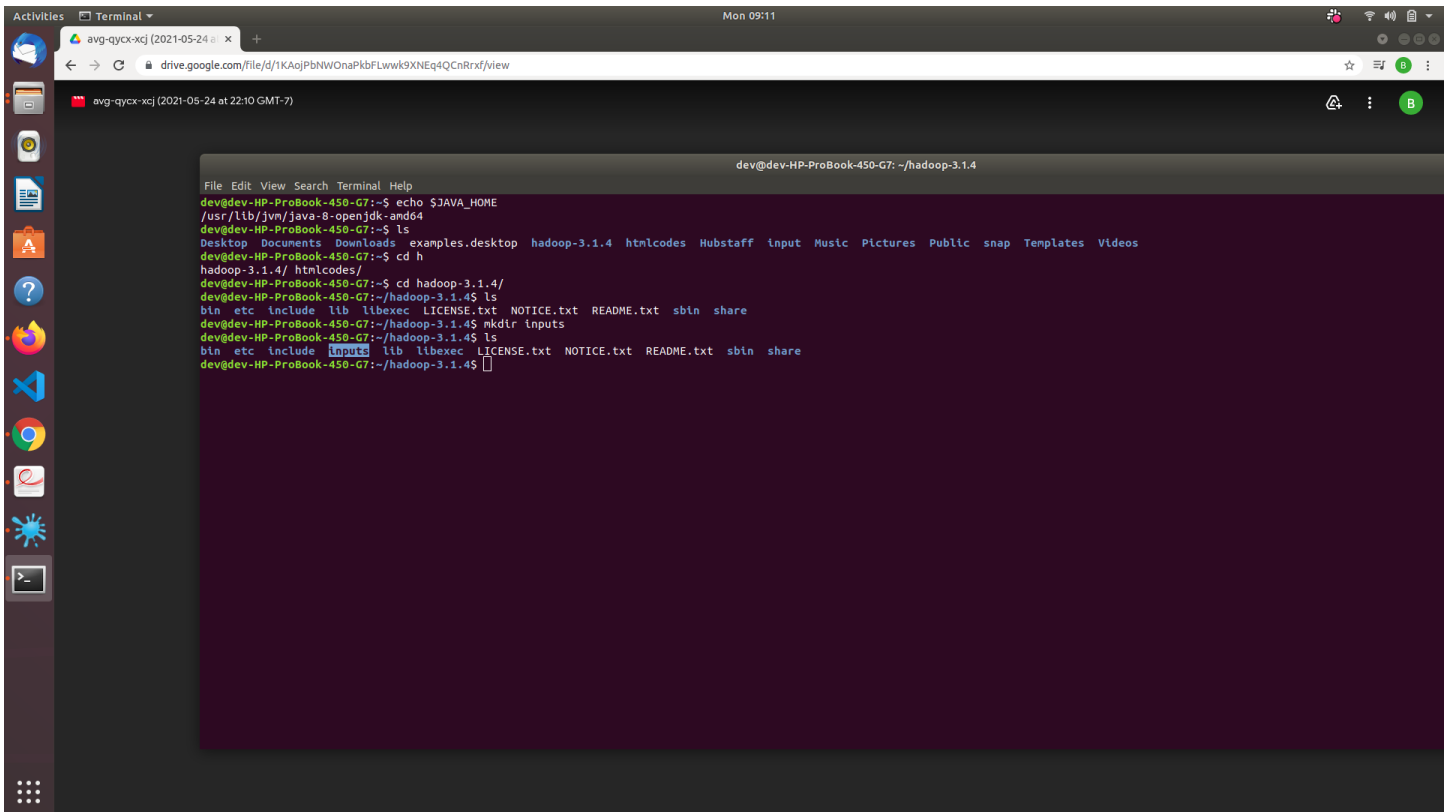
Cancel OK

Once the download is complete, extract the files to initiate the Hadoop installation:

```
tar xzf hadoop-3.2.1.tar.gz
```

The Hadoop binary files are now located within the `hadoop-3.2.1` directory.

[Live Chat](#)



```
2021-06-14 10:10:33,067 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2021-06-14 10:10:33,222 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2021-06-14 10:10:33,286 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2021-06-14 10:10:33,286 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2021-06-14 10:10:33,379 INFO input.FileInputFormat: Total input files to process : 0
2021-06-14 10:10:33,392 INFO mapreduce.JobSubmitter: number of splits:0
2021-06-14 10:10:33,484 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1671980215_0001
2021-06-14 10:10:33,484 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-06-14 10:10:33,585 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2021-06-14 10:10:33,585 INFO mapreduce.Job: Running job: job_local1671980215_0001
2021-06-14 10:10:33,586 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2021-06-14 10:10:33,592 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-06-14 10:10:33,592 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2021-06-14 10:10:33,593 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2021-06-14 10:10:33,618 INFO mapred.LocalJobRunner: Waiting for map tasks
2021-06-14 10:10:33,618 INFO mapred.LocalJobRunner: map task executor complete.
2021-06-14 10:10:33,621 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2021-06-14 10:10:33,622 INFO mapred.LocalJobRunner: Starting task: attempt_local1671980215_0001_r_000000_0
2021-06-14 10:10:33,640 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-06-14 10:10:33,640 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2021-06-14 10:10:33,652 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2021-06-14 10:10:33,654 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@1c343f0d
2021-06-14 10:10:33,655 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2021-06-14 10:10:33,666 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=3240257024, maxSingleShuffleLimit=810064256, mergeThreshold=2138569728, ioSortFactor=10, memoryThreshold=10
2021-06-14 10:10:33,668 INFO reduce.EventFetcher: attempt_local1671980215_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2021-06-14 10:10:33,673 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2021-06-14 10:10:33,676 INFO mapred.LocalJobRunner:
2021-06-14 10:10:33,676 INFO reduce.MergeManagerImpl: finalMerge called with 0 in-memory map-outputs and 0 on-disk map-outputs
2021-06-14 10:10:33,677 INFO reduce.MergeManagerImpl: Merging 0 files, 0 bytes from disk
2021-06-14 10:10:33,677 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2021-06-14 10:10:33,680 INFO mapred.Merger: Merging 0 sorted segments
2021-06-14 10:10:33,680 INFO mapred.Merger: Down to the last merge-pass, with 0 segments left of total size: 0 bytes
2021-06-14 10:10:33,681 INFO mapred.LocalJobRunner:
2021-06-14 10:10:33,707 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2021-06-14 10:10:33,718 INFO mapred.Task: Task:attempt_local1671980215_0001_r_000000_0 is done. And is in the process of committing
2021-06-14 10:10:33,719 INFO mapred.LocalJobRunner:
2021-06-14 10:10:33,719 INFO mapred.Task: Task attempt_local1671980215_0001_r_000000_0 is allowed to commit now
2021-06-14 10:10:33,721 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1671980215_0001_r_000000_0' to file:/home/hdooop/hadoop-3.2.1/grep-temp-1466437422
2021-06-14 10:10:33,721 INFO mapred.LocalJobRunner: reduce > reduce
2021-06-14 10:10:33,721 INFO mapred.Task: Task 'attempt_local1671980215_0001_r_000000_0' done
2021-06-14 10:10:34,612 INFO mapred.LocalJobRunner: Finishing task: attempt_local1671980215_0002_r_000000_0
2021-06-14 10:10:34,812 INFO mapred.LocalJobRunner: reduce task executor complete.
2021-06-14 10:10:35,710 INFO mapreduce.Job: Job job_local305787015_0002 running in uber mode : false
2021-06-14 10:10:35,711 INFO mapreduce.Job: map 100% reduce 100%
2021-06-14 10:10:35,711 INFO mapreduce.Job: Job job_local305787015_0002 completed successfully
2021-06-14 10:10:35,712 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=1266784
    FILE: Number of bytes written=3356368
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
  Map-Reduce Framework
    Map input records=0
    Map output records=0
    Map output bytes=0
    Map output materialized bytes=6
    Input split bytes=128
    Combine input records=0
    Combine output records=0
    Reduce input groups=0
    Reduce shuffle bytes=6
    Reduce input records=0
    Reduce output records=0
    Spilled Records=0
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=0
    Total committed heap usage (bytes)=790626304
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=98
  File Output Format Counters
    Bytes Written=8
```

```
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=2239
  Total time spent by all reduces in occupied slots (ms)=2504
  Total time spent by all map tasks (ms)=2239
  Total time spent by all reduce tasks (ms)=2504
  Total vcore-milliseconds taken by all map tasks=2239
  Total vcore-milliseconds taken by all reduce tasks=2504
  Total megabyte-milliseconds taken by all map tasks=2292736
  Total megabyte-milliseconds taken by all reduce tasks=2564096
Map-Reduce Framework
  Map input records=3
  Map output records=8
  Map output bytes=83
  Map output materialized bytes=80
  Input split bytes=110
  Combine input records=0
  Combine output records=0
  Reduce input groups=0
  Reduce shuffle bytes=80
  Reduce input records=6
  Reduce output records=6
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=154
  CPU time spent (ms)=1450
  Physical memory (bytes) snapshot=524750848
  Virtual memory (bytes) snapshot=5219407744
  Total committed heap usage (bytes)=456138560
  Peak Map Physical memory (bytes)=792489344
  Peak Map Virtual memory (bytes)=2006850048
  Peak Reduce Physical memory (bytes)=232341504
  Peak Reduce Virtual memory (bytes)=2612637696
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=51
```